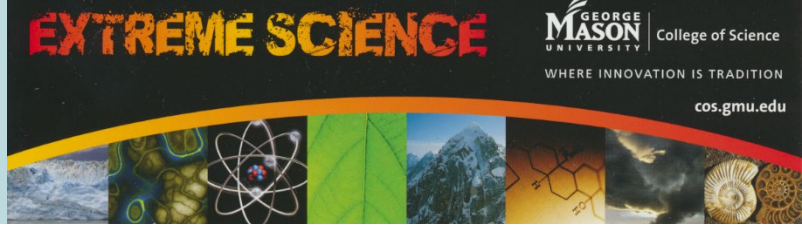

Lecture notes

CDS 302

Scientific Data and Databases

Fall Semester 2020

Lecture 1: Introduction

Lecture: Joe Boone

Who am I?

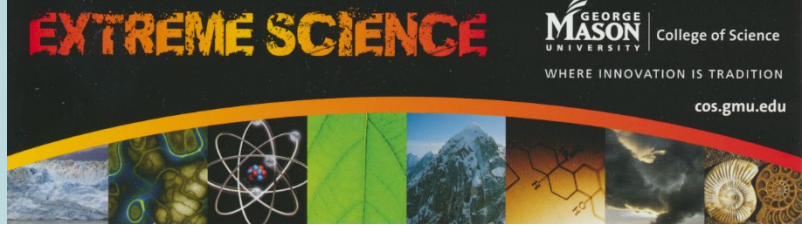

- Joe Boone – GMU Graduate Lecturer
  - jboone@gmu.edu
- Academic Background
  - BS and MS in Computer Science from GMU
  - Currently a Computational Science and Informatics Ph.D. Student at GMU
- Professional Career
  - 30+ years of Systems Development and Engineering
  - Satellite Telecommunication Applications
  - Geospatial Applications
  - Extensive Software Development Experience
  - Graduate Lecturer at GMU

Week 1 Topics

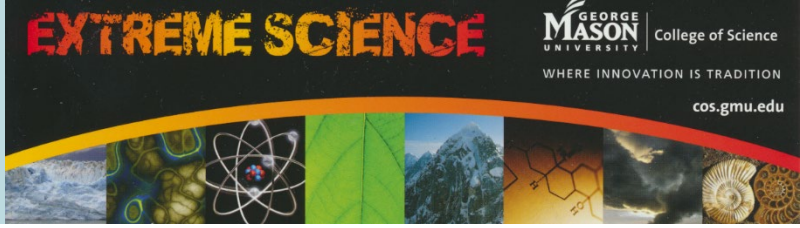

- Syllabus Review
- Tools
- Introduction to Database Systems
- Introduction to Scientific Writing
- Introduction to LaTeX and Overleaf (Part 1)
- Week 1 Assignments

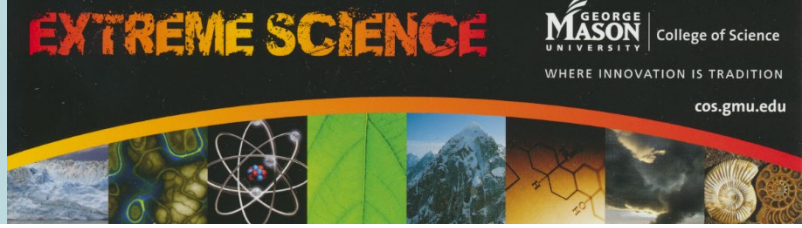

Syllabus Review

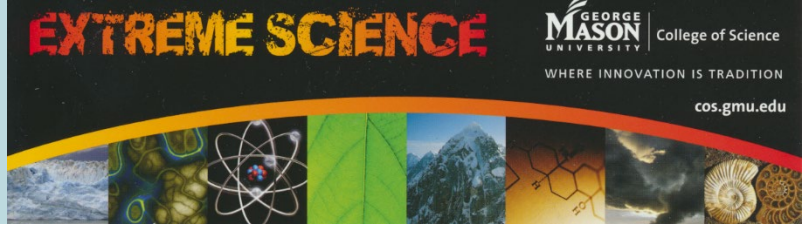
EXTREME SCIENCE
GEORGE MASON UNIVERSITY
College of Science
WHERE INNOVATION IS TRADITION
cos.gmu.edu

Tools

Software You Will Need

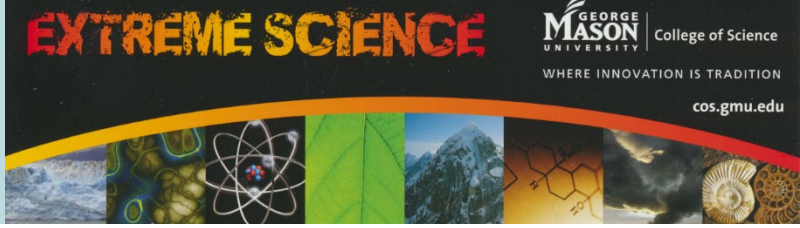

www.overleaf.com
www.sqlite.org/download.html
www.sqlitebrowser.org/dl/

www.anaconda.com/distribution/

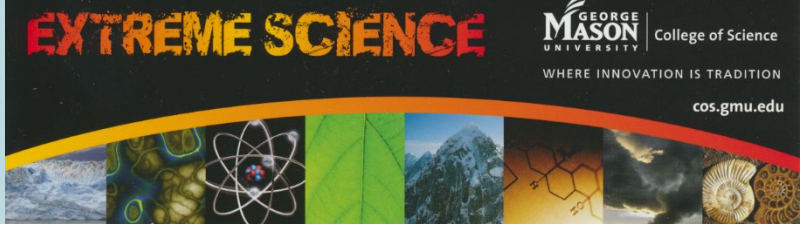

An Introduction to Database Systems

Some Basic Terminology

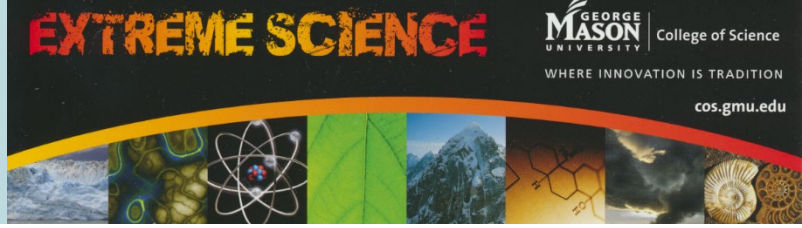

- **Database**
  - Database: Data collection, typically large and stored in secondary memory
  - Contains interrelated data on some enterprise

- **Database Management System (DBMS)**
  - Database Management System: a software package (i.e., collection of programs) designed to *define, manipulate, retrieve* and *manage* data in a database
  - Provides an environment that is convenient and efficient to use for multiple users and applications simultaneously

- **Database System**
  - DBMS + Database

Data Sizes Reference

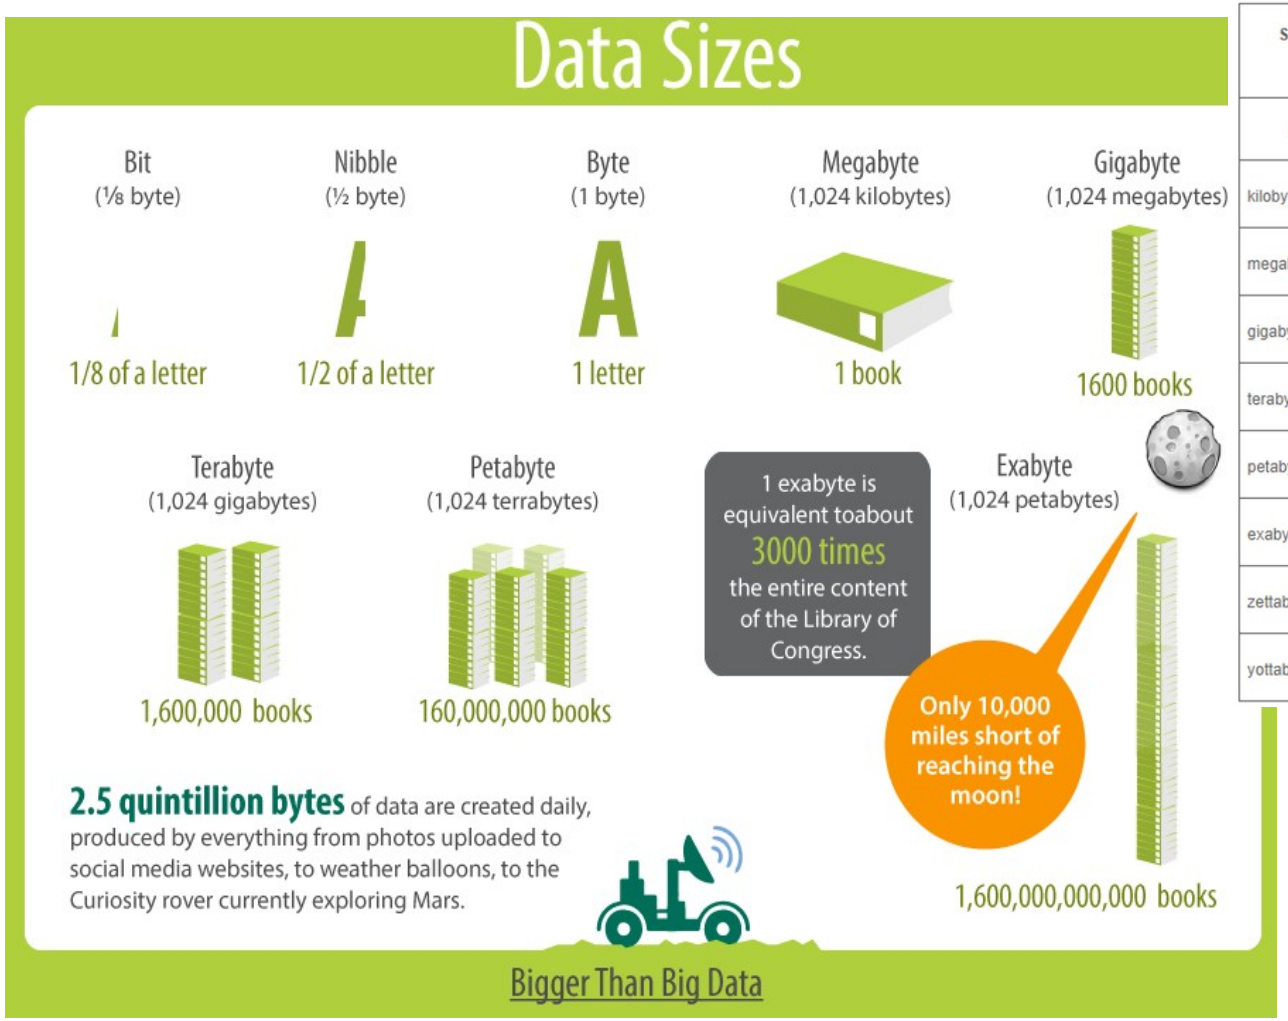

SI decimal prefixes		IEC binary prefixes		Percentage Difference IEC/SI	
Name	Value	Name	Value		
kilobyte (kB)	10^3	kibibyte (KiB)	2^{10}	2.4%	
megabyte (MB)	10^6	mebibyte (MiB)	2^{20}	4.9%	
gigabyte (GB)	10^9	gibibyte (GiB)	2^{30}	7.4%	
terabyte (TB)	10^{12}	tebibyte (TiB)	2^{40}	10.0%	
petabyte (PB)	10^{15}	pebibyte (PiB)	2^{50}	12.6%	
exabyte (EB)	10^{18}	exbibyte (EiB)	2^{60}	15.3%	
zettabyte (ZB)	10^{21}	zebibyte (ZiB)	2^{70}	18.1%	
yottabyte (YB)	10^{24}	yobibyte (YiB)	2^{80}	20.9%	

IEC and SI Size Notations

base2 Notation

```
IEC Notation   Size
-------------- ------------------------------------------------
KiB = kibibyte (2^10  / 1,024 bytes)
MiB = mebibyte (2^20  / 1,048,576 bytes)
GiB = gibibyte (2^30  / 1,073,741,824 bytes)
TiB = tebibyte (2^40  / 1,099,511,627,776 bytes)
PiB = pebibyte (2^50  / 1,125,899,906,842,624 bytes)
EiB = exbibyte (2^60  / 1,152,921,504,606,846,976 bytes)
ZiB = zebibyte (2^70  / 1,180,591,620,717,411,303,424 bytes)
YiB = yebibyte (2^80  / 1,208,925,819,614,629,174,706,176 bytes)
```

base10 Notation

```
SI Notation    Size
-------------- ------------------------------------------------
KB = kilobyte  (10^3  / 1,000 bytes)
MB = megabyte  (10^6  / 1,000,000 bytes)
GB = gigabyte  (10^9  / 1,000,000,000 bytes)
TB = terabyte  (10^12 / 1,000,000,000,000 bytes)
PB = petabyte  (10^15 / 1,000,000,000,000,000 bytes)
EB = exabyte   (10^18 / 1,000,000,000,000,000,000 bytes)
ZB = zettabyte (10^21 / 1,000,000,000,000,000,000,000 bytes)
YB = yottabyte (10^24 / 1,000,000,000,000,000,000,000,000 bytes)
```

Why do we need Database Systems? Data is Everywhere

- Huge flood of data
  - Modern technology
  - New user mentality
  - 2.5 Exabytes of new data every day
- New applications
- Innovative research
- Economic Boost
  - "$600 billion potential annual consumer surplus from using personal location data" **[1]**

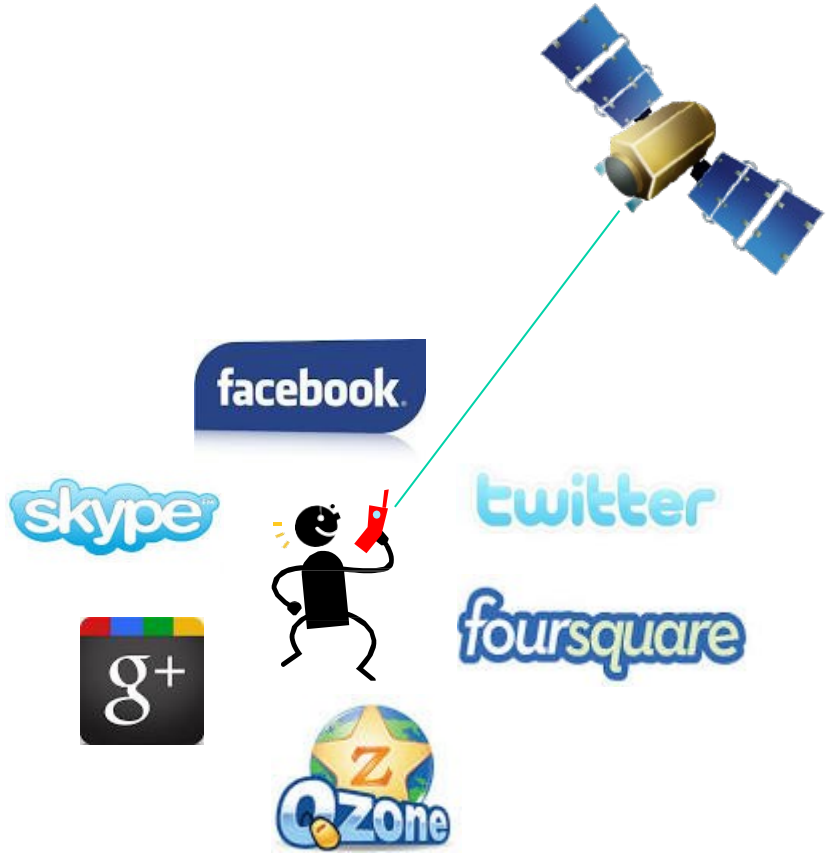

[1] McKinsey Global Institute. Big data: The next frontier for innovation, competition, and productivity. June 2011.

Why do we need Database Systems? Data Production is Accelerating

Quoted in 2010…

Why do we need Database Systems?
Big Data = $$$

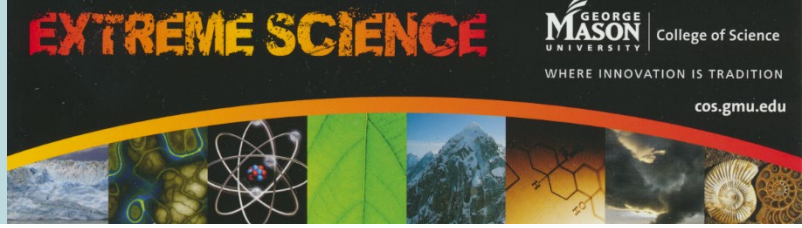

Opportunity…

The 3-*ish* V's of Big data…

- Volume
- Velocity
- Variety

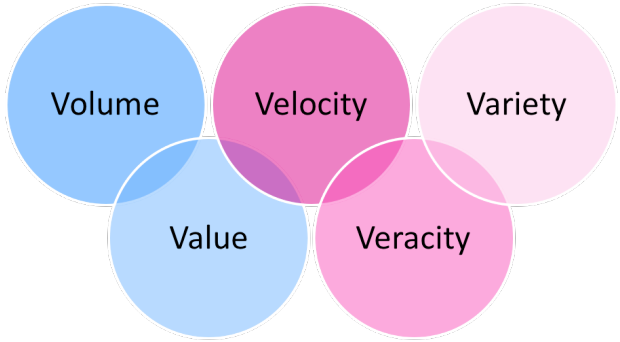

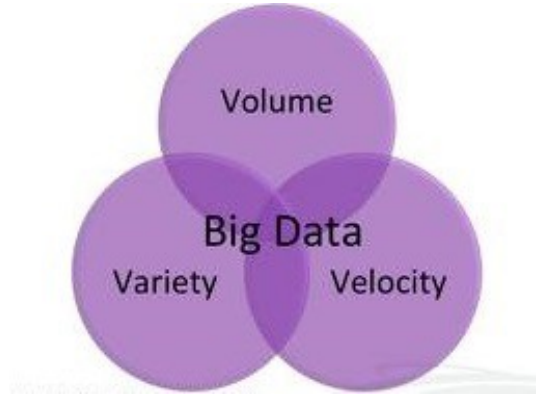

- *Veracity*
- *Variability*
- *Visualization*
- *Value*

…

Why do we need Database Systems?
Big Data Hype

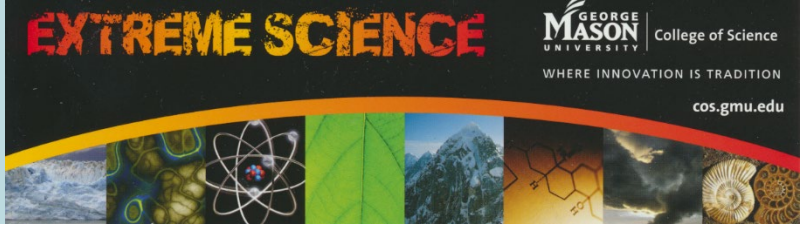

- Big Data analytics is a fancy new word for Knowledge Discovery in Databases!

- KDD has focused on large data for decades:

41st International Conference on
VERY LARGE DATA BASES
Hilton Waikoloa Hotel • Kohala Coast, Hawai'i
August 31 - September 4, 2015

  - **VL**DB since 1975

- **Big Data is not new**

Why do we need Database Systems? Data Collection

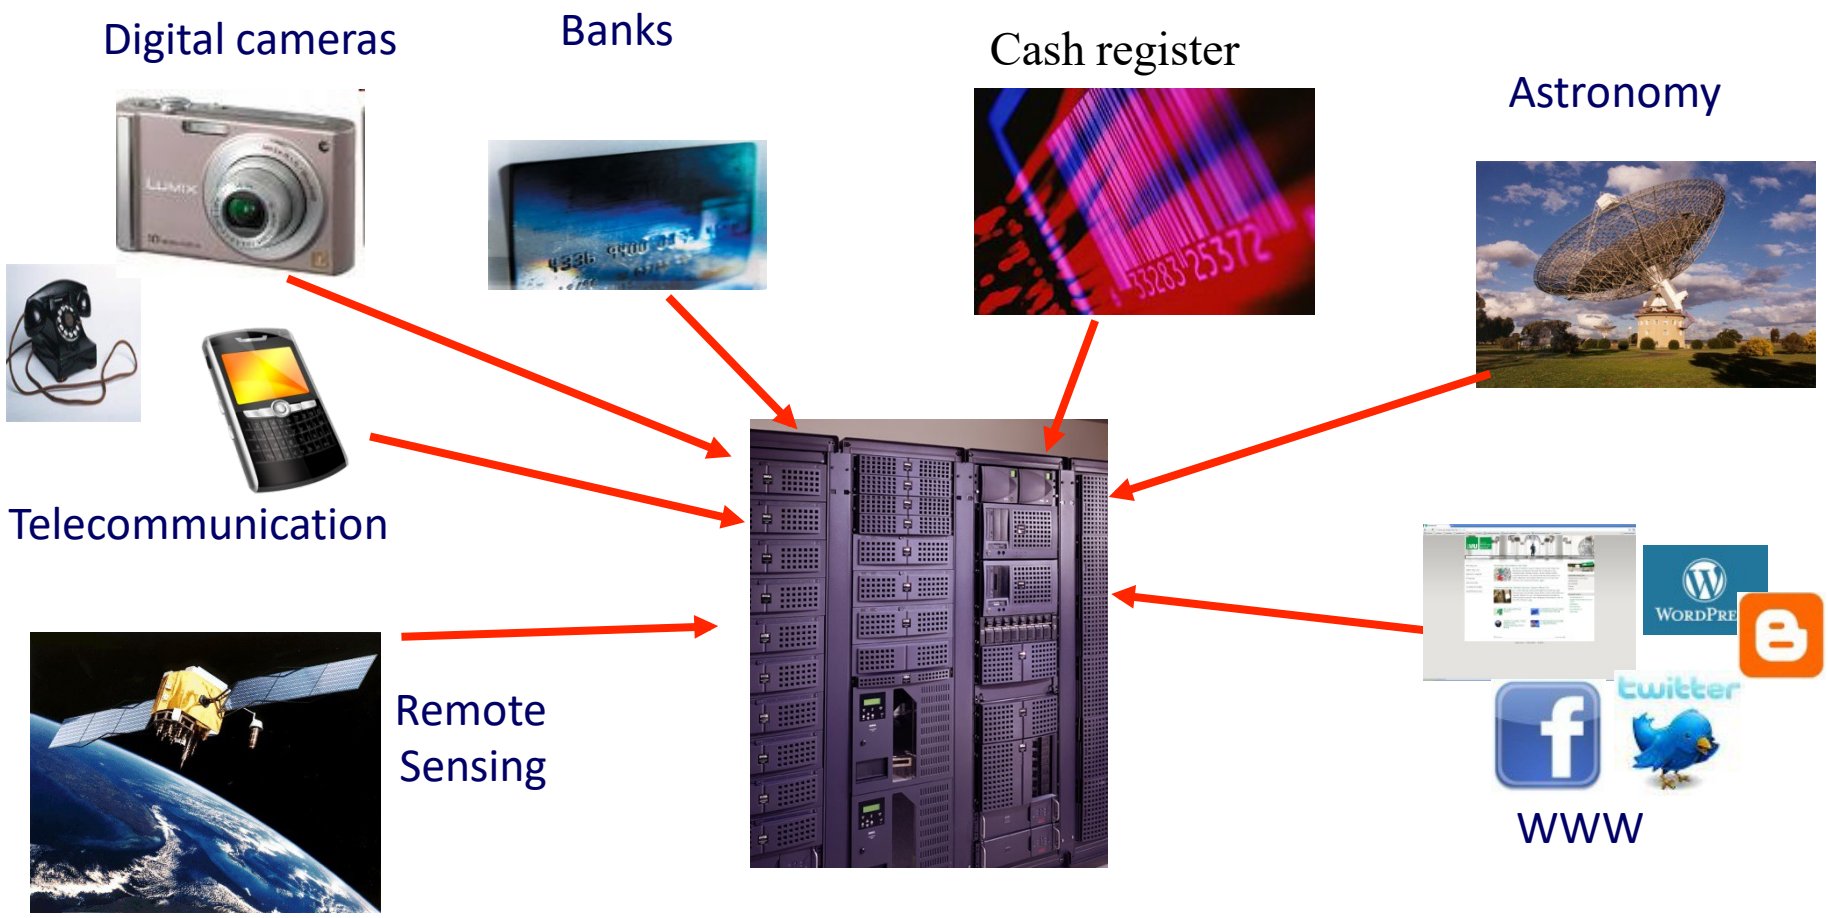

- Huge amounts of data are collected nowadays from different application domains
- Is not feasible to analyze all these data manually

Why do we need Database Systems?
Who uses databases?

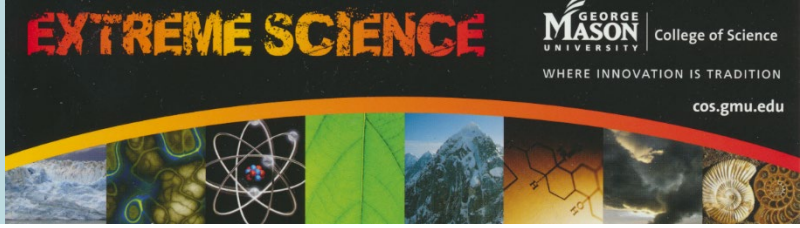

Zillow

tripadvisor

amazon

ebay

Everyone…

Why do we need Database Systems?

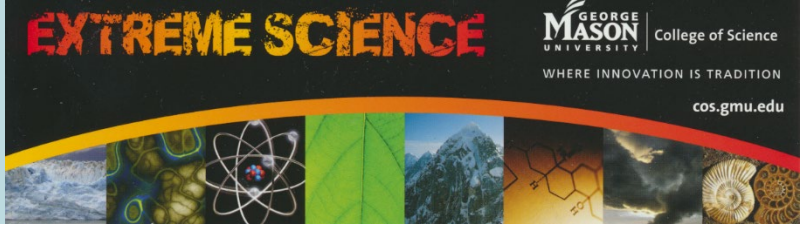

- What problems do database systems solve:
  - **Data redundancy and inconsistency**
    - Duplication of data, different values for the same variable
  - **Difficulty in accessing the data**
  - **Data isolation**
    - Different files and formats complexity is hidden from users (abstracted)
  - **Data integrity**
    - Data is not logically consistent and rules to enforce are encoded in program code
    - Hard to add/modify constraints
  - **Atomicity**
    - Failures in the middle of a transaction leave the system in an inconsistent state
  - **Concurrent access by multiple users**
  - **Managing security**
    - Need to manage who has access to what data
  - **Proprietary systems lead to vendor "lock-in" and other problems**
  - **Performance and scalability**

Terminology: *ACID Transaction,* Atomicity, Consistency, Isolation, Durability

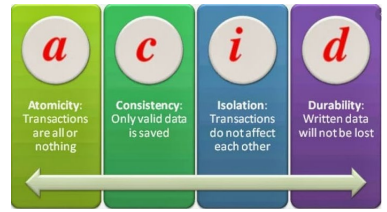

History of Database Systems

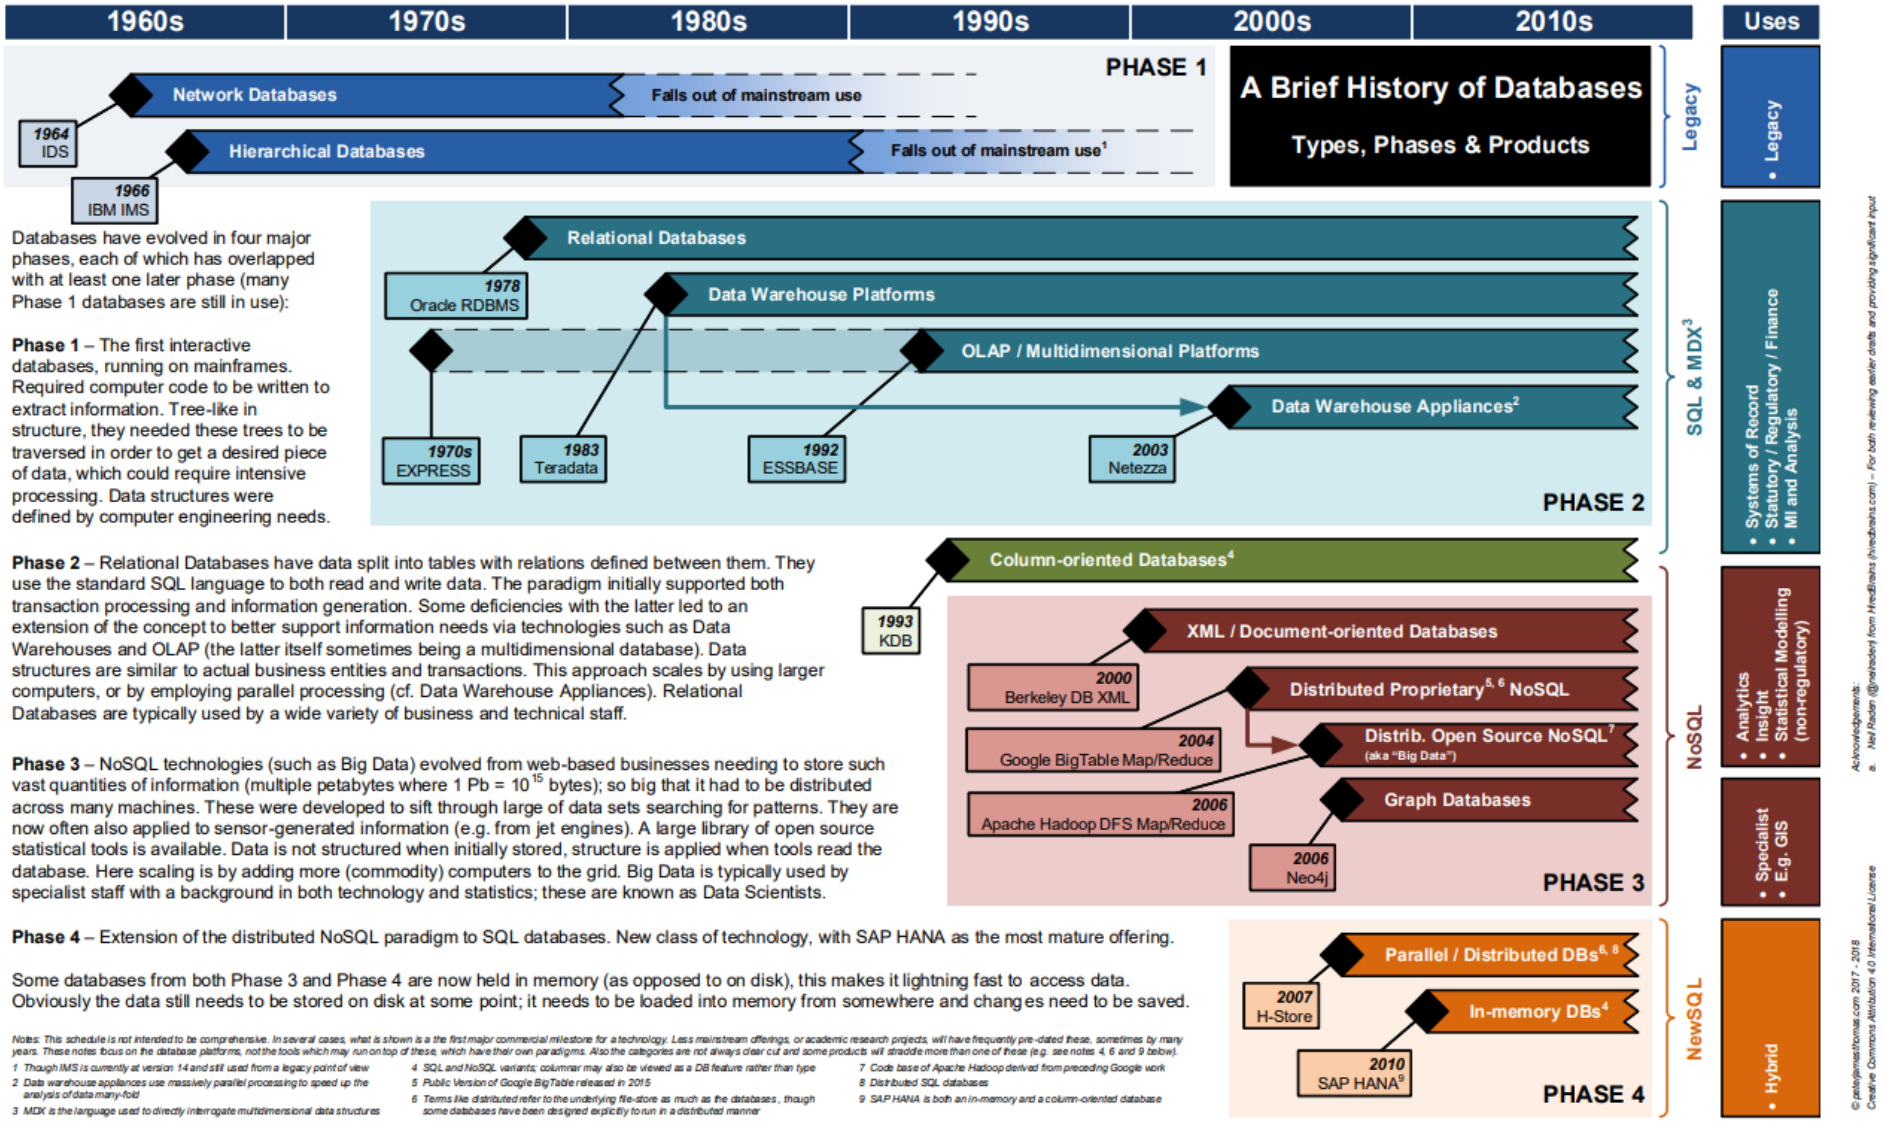

Slide from https://peterjamesthomas.com/2018/02/06/a-brief-history-of-databases/

Database Systems by Popularity

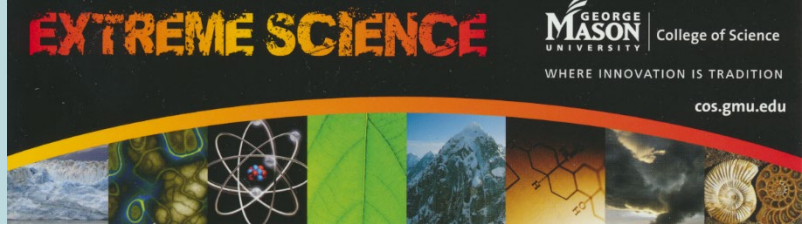

359 systems in ranking, August 2020

Rank Aug 2020	Rank Jul 2020	Rank Aug 2019	DBMS	Database Model	Score Aug 2020	Score Jul 2020	Score Aug 2019
1.	1.	1.	Oracle	Relational, Multi-model	1355.16	+14.90	+15.68
2.	2.	2.	MySQL	Relational, Multi-model	1261.57	-6.93	+7.89
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model	1075.87	+16.15	-17.30
4.	4.	4.	PostgreSQL	Relational, Multi-model	536.77	+9.76	+55.43
5.	5.	5.	MongoDB	Document, Multi-model	443.56	+0.08	+38.99
6.	6.	6.	IBM Db2	Relational, Multi-model	162.45	-0.72	-10.50
7.	↑ 8.	↑ 8.	Redis	Key-value, Multi-model	152.87	+2.83	+8.79
8.	↓ 7.	↓ 7.	Elasticsearch	Search engine, Multi-model	152.32	+0.73	+3.23
9.	9.	↑ 11.	SQLite	Relational	126.82	-0.64	+4.10
10.	↑ 11.	↓ 9.	Microsoft Access	Relational	119.86	+3.32	-15.47

https://db-engines.com/en/ranking

Relational databases are still dominant and provide a foundation to build on…

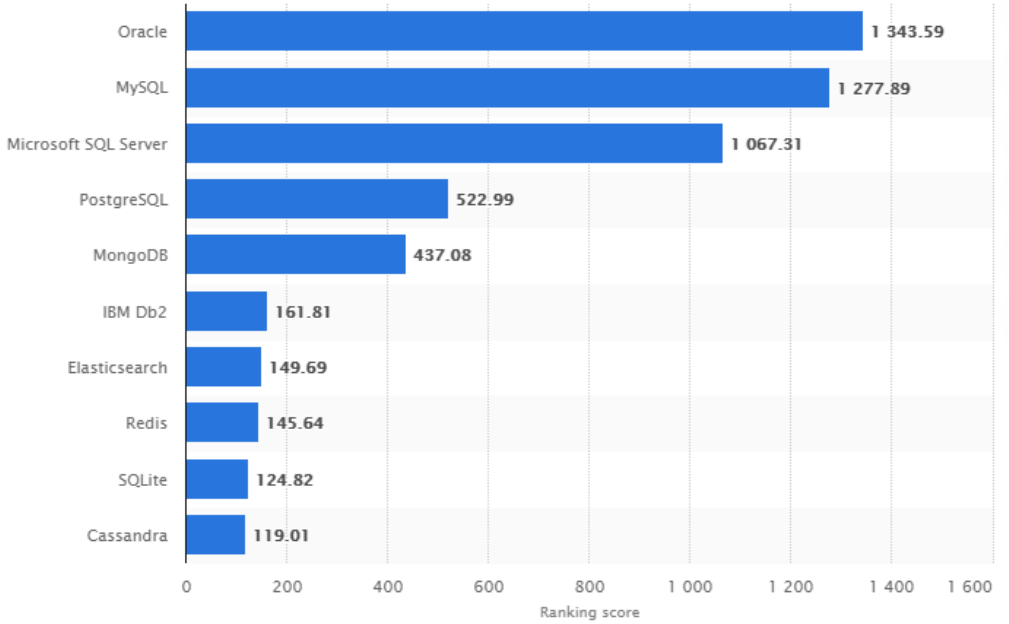

https://www.statista.com/statistics/809750/worldwide-popularity-ranking-database-management-systems/

Why Relational DBMS?

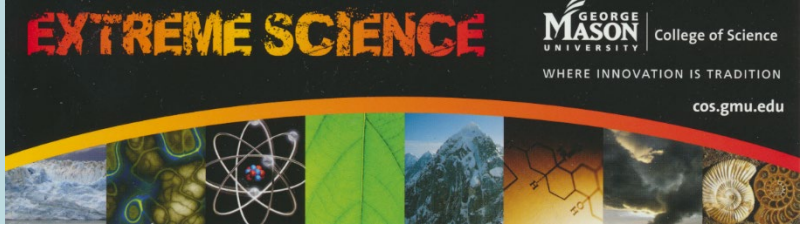

- **In our course we will be focusing on Relational Database Systems:**
  - **Data Modeling**
    - Redundancy Control
    - Referential Integrity / Consistency Constraints
  - **Efficient Query Processing**
    - Indexing
    - Optimization
  - **Operating Accuracy**
    - Error recovery - Atomicity
    - Concurrent access by multiple users
  - **Security Issues**
    - Control access rights
  - **Standards Based**
    - Structured Query Language(SQL)is an ANSI and ISO standard
    - Kind of…
  - **Large Legacy Base**
    - Lots and lots of relational databases are out there…

Why not Relational DBMS?

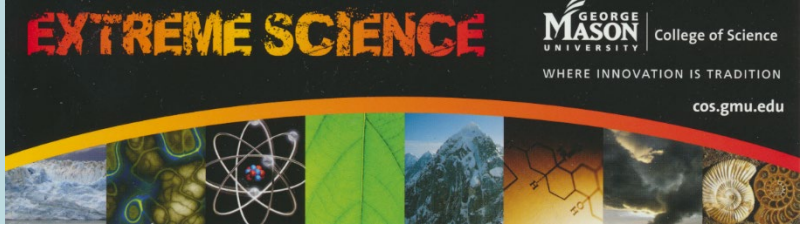

- Relational DBMS are designed for structured design and development
  - Rigidly enforced rules for data integrity
  - This is not applicable for all situations

- In modern situations, valuable data is often generated in a relatively unstructured format:
  - Examples: text, log files, documents, BLOBS(binary large objects) like pictures, videos, audio, etc…
  - While these can sometimes be held within a relational database – these data do not neatly fall into the relational model
  - The system being designed not require inherent Relational DBMS features: e.g. Atomicity, data integrity checking, or may have asymmetric requirements between storing and retrieving functions

Types of Database Systems?

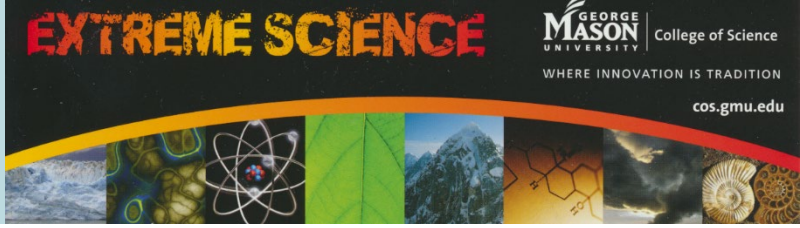

- **Relational (Classic Relational)**
  - Oracle, IBM/DB2, MS SQL Server
  - PostgreSQL, MySQL, SQLite
- **Non-Relational (NoSQL "Not Only SQL")**
  - Key-Value Store (e.g. AWS DynamoDB, redis)
  - Document Store (e.g. MongoDB)
  - Wide-Column Stores (examples: Bigtable, Cassandra, Apache Hbase)
  - Graph Database (e.g. Neo4j)
- **Object-Oriented Databases**
  - Object-Oriented Store (e.g. InterSystems Cache)
- **Geographic Databases**
  - Vectors/Raster Store (e.g. ESRI Geodatabases / SpatiaLite)
- **First Databases…(1960s)**
  - Network (IDS, Many-to-many relationship)
  - Hierarchical (IBM IMS, Parent-child relationship)
- **DBMS as a Service (Not really a type of database…)**
  - Google
  - AWS

DBMS Architecture (briefly…)

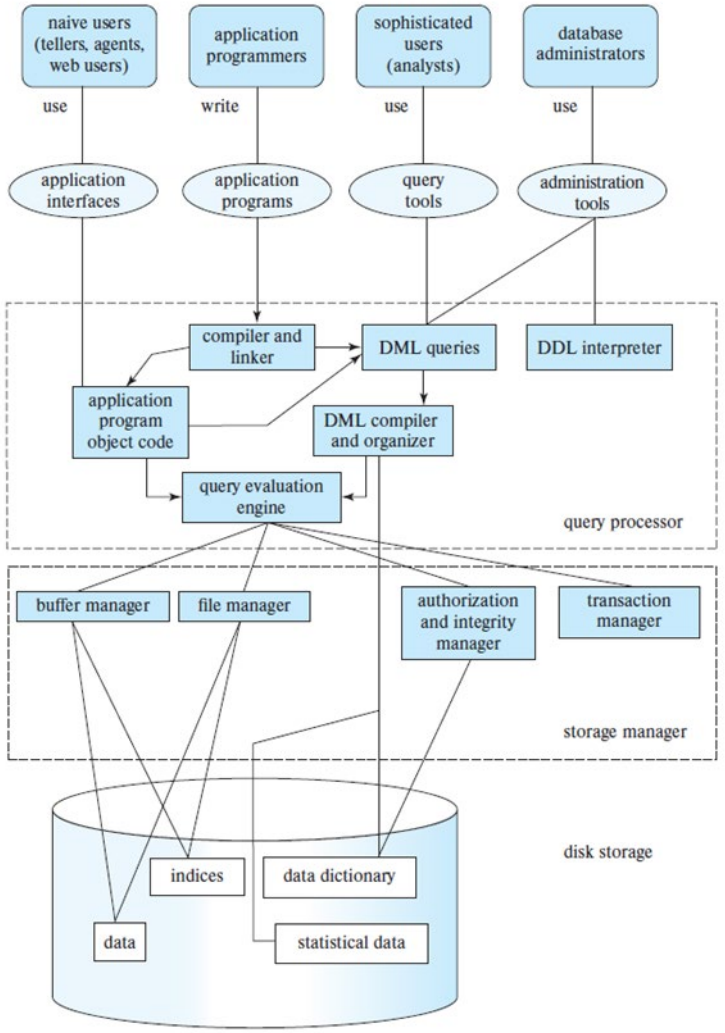

Figure 1.3 System structure.

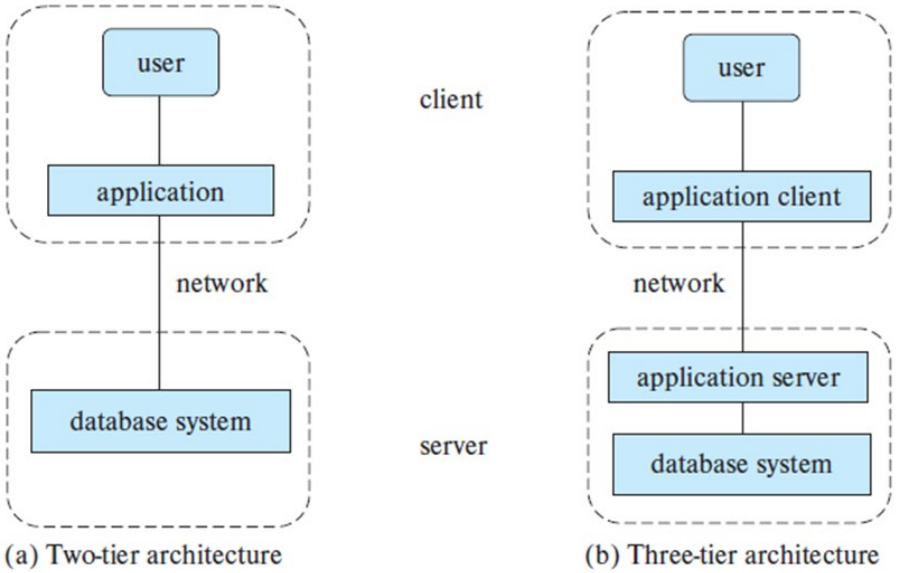

Figure 1.4 Two-tier and three-tier architectures.

Why do I want to learn it?

DATA

Data Scientist: The Sexiest Job of the 21st Century [1]

by Thomas H. Davenport and D.J. Patil

If this means having rare qualities that are much in demand, data scientists are already there. They are difficult and expensive to hire and, given the very competitive market for their services, difficult to retain.

[1] Harvard Business Review. Data Scientist: The Sexiest Job of the 21st Century. October 2012.

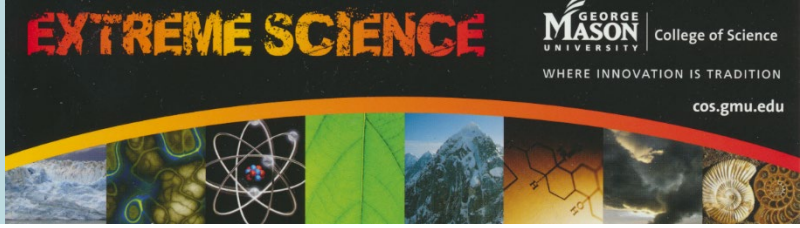

An Introduction to Scientific Writing

Why read journal articles?

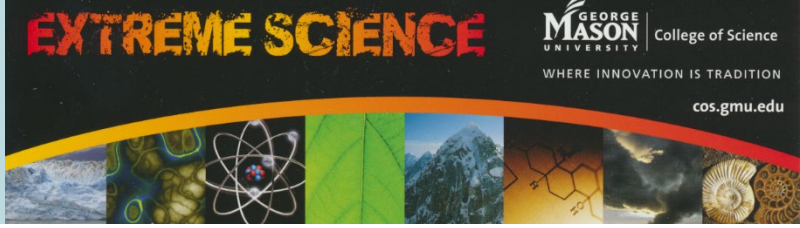

- To stay current with the progress in a field of study
- To find the solution for a specific problem
  - Test / Methods
- To understand the fundamental background in an area of study
- To get an idea for carrying out further research
- You have been assigned to review the article by a Professor or Journal Editor
- To support, refine, refute your scientific beliefs or views
- To impress others…

Adapted from: How to read clinical journals: I. why to read them and how to start reading them critically. Can Med Assoc J. 1981 Mar 1; 124(5):555-8; Durbin CG., Jr How to read a scientific research paper. Respir Care. 2009;54:1366–71.

Types of articles published in a scientific journal

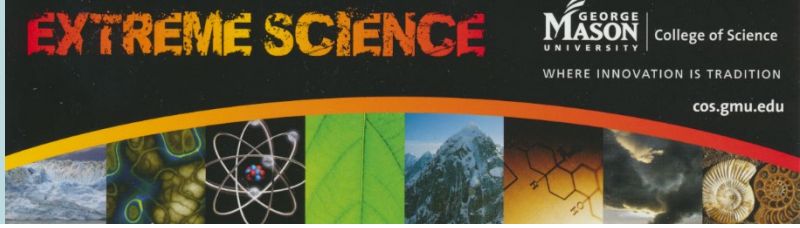

- **Primary literature**
  - "Core" of scientific publications
  - Present findings on new scientific discoveries
  - Describe earlier work to acknowledge it and place new findings in the proper perspective
    - Original research articles
    - Surveys
    - Case report/case series
    - Conference proceedings and abstracts
    - Editorial
    - Correspondence/letters to the Editor
- **Secondary literature**
  - Original research information is reviewed
    - Narrative reviews
    - Systematic reviews
    - Meta-analysis
    - Book reviews
    - Guidelines
    - Commentary

Structure of a journal article

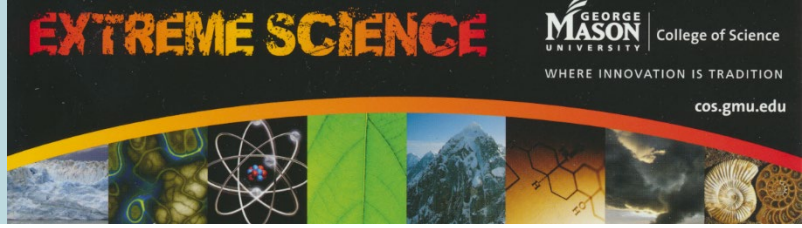

- **Title:** Topic and information about the authors.
- **Abstract:** Brief overview of the article.
- **Introduction:** Background information, gap in research and statement of the research hypothesis. Include motivation – why is this important.
- **Methods:** Details on how the study was conducted, procedures followed, instruments used, and variables measured. Must be systematic.
- **Results/Experimental Evaluation:** All the data of the study along with figures, tables and/or graphs.
- **Discussion:** The interpretation of the results and implications of the study. Were the objectives met? Limitations and future work.
- **Conclusion:** What does all this mean?
- **References/Bibliography:** Citations of sources from where the information was retained.

Reading (Filtering) process

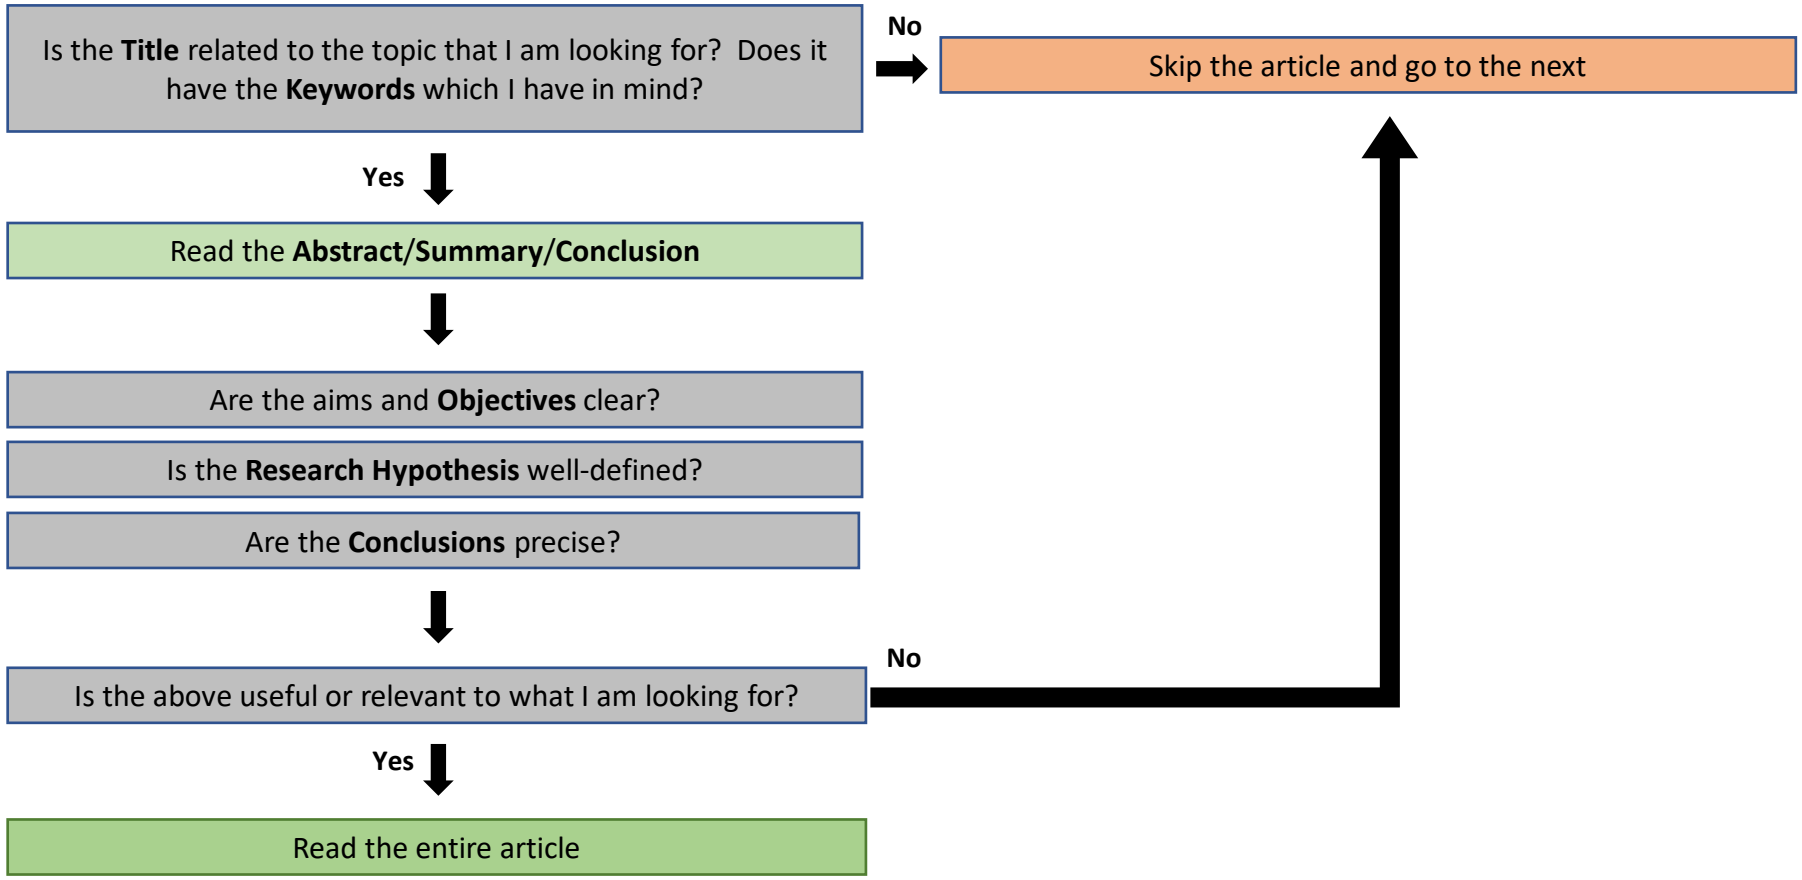

Research Questionnaire

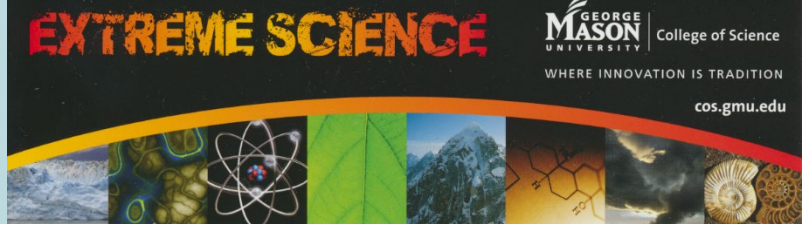

Research Questionnaire

Overall

1. What was the article type?
2. What was the title?
3. Who were the authors?

Introduction

4. What was the research problem?
5. Was there any mention of previous studies on this topic?
6. Why was this study performed (the rationale)?
7. What were the aims and objectives of the study?
8. What was the study (research) hypothesis?

Materials and Methods

9. How did the researcher attempt to answer the research question?
10. How was the sampling done?
11. How were they grouped (categorized)?
12. What were the inclusion criteria?
13. What were the exclusion criteria?
14. What procedures were followed?
15. Which variables were measured?
16. What equipment/instruments were used for data collection? Were they appropriate?
17. What statistical methods/tests were employed? Were they apt for evaluation?

Results

18. What were the key findings?
19. Were all the subjects present in the beginning of the study accounted for at the end of the study?
20. Were the results reliable?
21. Were the results valid?
22. Which results were statistically significant?
23. Which results were statistically non-significant?
24. Were the tables/graphs easy to comprehend?

Discussion

25. Did the results answer the research question?
26. What were the author's interpretations of the data?
27. Was the analysis of the data relevant to the research question?
28. How were those results different/similar when compared to other studies?
29. What were the strengths of the study?
30. What were the limitations of the study?
31. Were there any extrapolations of the findings beyond the range of data?

Conclusions

32. What were the conclusions?
33. Were the author's conclusions based upon reported data and analysis?
34. Were the conclusions reasonable and logical?
35. Will the results be useful in practice or for further research?
36. Was the study worth doing?
37. Does the read have any questions unanswered by the article?

References

38. Were the references cited according to the journal's requirements?
39. Were all the citations correct?
40. Were all the references cited in the text?

It takes time…

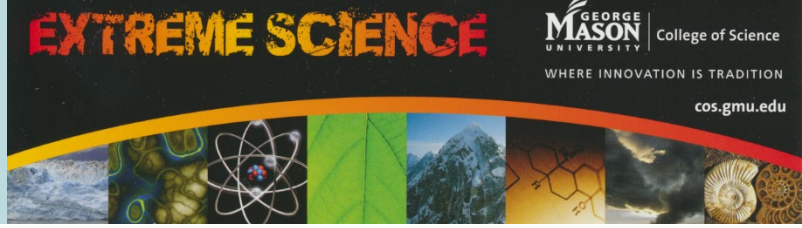

- It the not the same as reading a novel or a blog
- It’s a skill
- It gets better with practice
- It gets better as you become more familiar with the research area
- The first paper may take some time
- You may have to look at other resources to understand some of the paper’s content
- Be patient....you’ll get there!

Additional resources

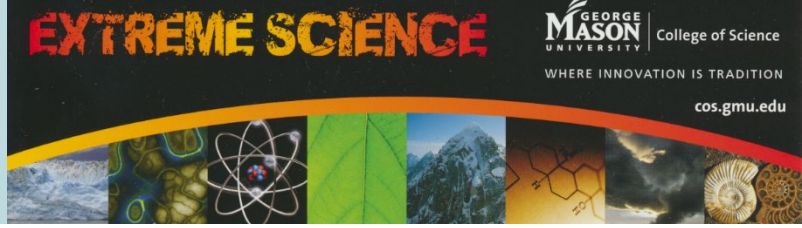

- Google Search
- Google Scholar
- GMU Library resources (library.gmu.edu)
- Academia.edu
- ResearchGate
- Reddit Scholar
- Email scholars if you can’t get their articles freely

- Citations managers like Zotero or Mendeley
  - Help you keep your research materials organized and in one place
  - Help you generate a bibliography

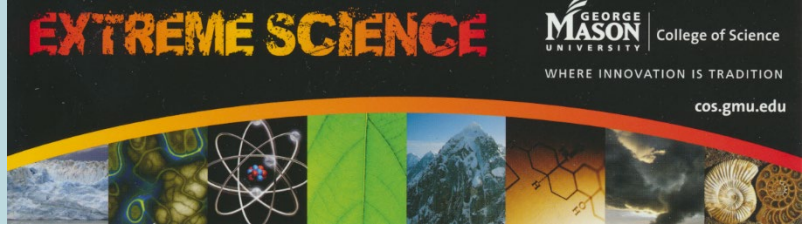

Introduction to LaTeX

LaTeX

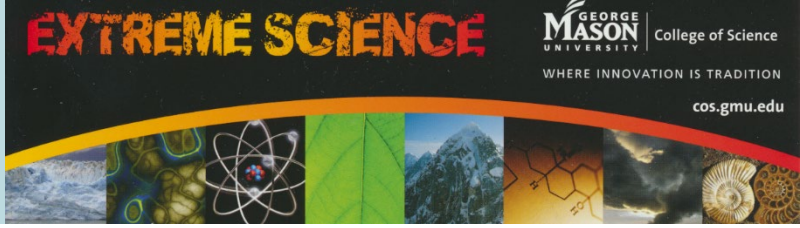

- LaTeX
  - LaTeX is “Lamport” + “TeX”
  - Leslie Lamport
    - Computer Scientist, Distributed Systems (2013 Turing Award)
    - Initial Developer / Inventor of LaTeX (1983)
    - LaTeX is a tool for document preparation built on top of the typesetting system TeX
    - It is the standard for scientific journal articles
    - Pronounced (“la” or “lay”) + “tech”
  - TeX
    - Late 1970’s
    - Donald Knuth, Computer Scientist (Turing Award 1974, many awards…)
    - Author of *The Art of Computer Programming* (classic CS text)
    - Typesetting engine that drives LaTeX and other higher-level packages

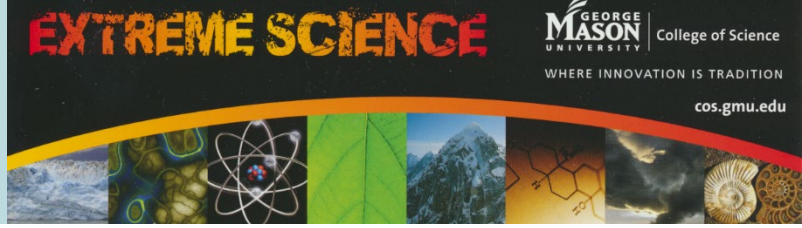

Assignments

Assignments Week 1

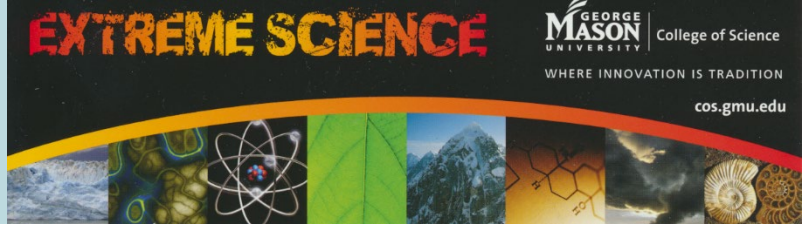

- Setup a free account with Overleaf (www.overleaf.com)
  - Follow along with the recorded lectures covering LaTeX
  - Explore LaTeX on your own…(Next week we will cover Part 2)

- Suggested reading: Silberschatz et. al., *Database System Concepts, McGraw-Hill,* Chapter 1

Other Resources

- *Introduction to Databases* class by Jennifer Widom, Stanford
  - http://www.db-class.org/course/auth/welcome

- LaTeX tutorials
  - https://www.overleaf.com/learn/latex/Tutorials
  - https://www.latex-tutorial.com/tutorials/

- List of LaTeX Math Symbols
  - https://www.caam.rice.edu/~heinken/latex/symbols.pdf